


BREADTH-FIRST SEARCH (BFS)

Intuition: BFS(vertex v)


To start, all vertices are *unmarked*.

- ▶ *Start* at v . *Visit* v and mark as *visited*.
- ▶ *Visit* every *unmarked* neighbour u_i of v and mark each u_i as *visited*.
- ▶ *Mark* v *finished*.
- ▶ *Recurse* on each vertex marked as *visited* in the order they were visited.



BFS

Q: What *information* about the graph can a *BFS* be used to find?



- ▶ the *shortest path* from v to any other vertex u and this distance $d(v)$.
- ▶ Whether the graph is *connected*.
- ▶ Number of *connected components*.

Q: What does the *BFS* construct?

A: a *BFS tree* that visits every node *connected* to v , we call this a *spanning-tree*.

IMPLEMENTING BFS

Q: What is an appropriate **ADT** to **implement** a **BFS** given an **adjacency list** representation of a graph?

A. A FIFO (first in, first out) queue. which has the operations:

- ▶ `ENQUEUE (Q, v)`,
- ▶ `DEQUEUE (Q)`,
- ▶ `ISEMPTY (Q)`

Q: What **information** will we need to store along the way for each v ?

- ▶ the current node u and it's **state** (ie, **visited**, **not visited**, **finished**)
- ▶ the predecessor $p[u]$
- ▶ possibly the **distance** $d[u]$ from v if needed
- ▶ $o[u]$ the **order** of discovery

THE BFS ALGORITHM

```
BFS(G=(V,E), v):      # Start BFS on G at vertex v
    for u in V: # Initialize arrays
        state[u]=not_visited; o[u]=-1; d[u]=infinity; p[u]=NIL
    new queue Q
    i = 1
    state[v] = visited          green = visited
    o[v] = i; d[v] = 0; p[v] = NIL    black = not visited
    ENQUEUE(Q,v)                red = finished
    while not ISEMPTY(Q):
        u = DEQUEUE(Q)
        for each edge (u,w) in E:
            if (state[w] == not_visited):
                state[w] = visited
                i += 1
                o[w] = i
                d[w] = d[u] + 1
                p[w] = u
                ENQUEUE(Q,w)
        state[u] = finished
```

COMPLEXITY OF BFS(G, v)

Q: How many times is each node **ENQUEUE**ed?

A. At *most once*, when it is **not visited**, at which point it is marked *visited*.

⇒ the *adjacency list of each node* is traversed at most once.

⇒ so that the total running time of **BFS** is $O(n + m)$ or *linear* in the size of the adjacency list.

NOTES:

- ▶ BFS will visit only those vertices that are *reachable* from v .
- ▶ If the graph is *connected (in the undirected case)* or *strongly-connected (in the directed case)*, then this will be all the vertices.
- ▶ If not, then we may have to call **BFS** on more than *one start vertex* in order to see the whole graph.

Exercise Prove that $d[u]$ really does represent the *length* of the *shortest path* (in terms of number of edges) from v to u .